




— Mem BW roof (272 GB/s)

— INT8 tensor core (~130 TOPS)

... dp4a CUDA core (~13 TOPS)

▲ S1: FP32

▼ S2: Naive INT8

■ S3: Warp-Opt

◆ S4: +Fused Ops

+ S5: +Fused Attn

★ S6: dp4a W8A8