


Memory Access Pattern: Uncoalesced vs. Coalesced GEMV

Stage 2: Naive INT8 (Uncoalesced Access)

32 cache-line fetches
(1 byte used each)

Warp (32 threads)

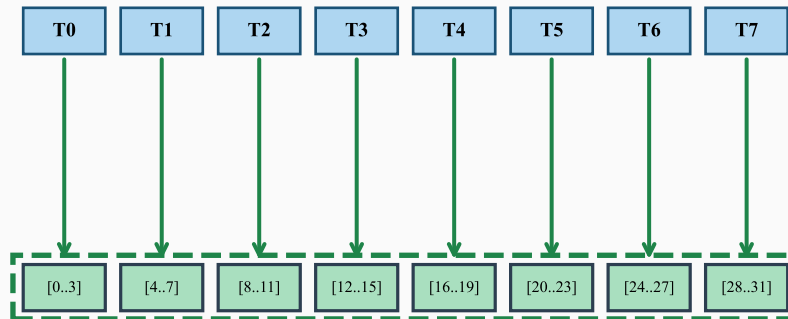


19.2 GB/s | L2 hit 2.6%
← stride = 2048 bytes →

Stage 3: Warp-Opt INT8 (Coalesced Access)

1 cache-line fetch
(128 bytes, all used)

Warp (32 threads)



85.6 GB/s | 5.7× speedup
← 128-byte coalesced transaction →